





Bayesian Phase Unwrapping with Factor Graphs

Eric Jonas

6.556 Final Project


May 12, 2009

Phase Unwrapping for MR




We know phase information varies smoothly except at boundaries

Markov Random Fields



Markov random field,
undirected graphical model,
etc.

Factor Graphs



- Factor Graphs [3] express the same concepts as MRFs but make the **factors** explicit.


Factor Graph Probability

$$P(x_1, x_2, x_3, x_4) = f_1(x_1, x_2) \cdot f_2(x_2, x_3) \cdot f_3(x_3, x_4) \cdot f_4(x_4, x_1) \quad (1)$$

Ising Model: The original lattice MRF

Imagine you want to simulate a spin system
Statistical physics
people love doing this

Factor Graphs for Low-Level Vision



Properties of Image Factor Graphs [1]:

- Use observed state (data) to infer hidden state
- Have lattice structure like the ising
- Large number of vertices ($O(n)$ in number of pixels)
- $O(1)$ per-vertex connectivity
- Typically have homogeneous factors

Bayesian Factor Graphs

Bayes Rule

$$P(X|Y) = \frac{P(Y|X)P(X)}{\sum_x P(Y|X)P(X)}$$


- $Y = y_{(i,j)}$: Observed Nodes
- $X = x_{(bi,j)}$: Hidden state we wish to estimate
- $P(Y|X)$: measurement model ("likelihood")
- $P(X)$: prior
-

Bayesian Factor Graphs For Low-Level vision

Our factor graph gives us $P(X, Y)$

MRFs for Phase: Frey's approach

In Ying and Frey's model [8] they formulate 2-D phase unwrapping as a low-level vision MRF problem.



- $Y_{(i,j)} \in \mathbb{R}$
- $x_{(i,j)} \in [0, 2\pi)$

Observation Potential

$$f_o = \delta((Y_{(i,j)} \bmod 2\pi) - X_{(i,j)})$$

Latent potential

$$f_l(X_1, X_2) = (X_1 - X_2)^2$$

discrete latent state, uniform factors

discrete latent state, unique factors

My formulation

Inference in MRFs

- The MRF tells us how to compute $P(X, Y)$. Bayes rule tells us how to compute $P(X|Y)$. But the sum is awful.
- But it's easy to compute $P^*(Y|X)$

Two generic approaches:

- draw samples from $p(x|D)$ to empirically estimate
- optimize to find MAP solution

Inference in MRFs

- The MRF tells us how to compute $P(X, Y)$. Bayes rule tells us how to compute $P(X|Y)$. But the sum is awful.
- But it's easy to compute $P^*(Y|X)$

Two generic approaches:

- draw samples from $p(x|D)$ to empirically estimate
- optimize to find MAP solution

We focus on sampling (Why?)

Markov-Chain Monte Carlo

Markov Property: next state only depends on current state

$$p(x_{t+1}|x_{1:t}) = p(x_{t+1}|x_t)$$

Ergodic markov chains have stationary distributions

Set up a state space so that the asymptotic limit is the target distribution

Used in situations where you want to sample from $\pi(x)$ but only can compute $\pi^*(x)$

Metropolis Hastings

Consider a proposal distribution, $q(x \rightarrow x^*)$ We draw x^* as a *new target state* from this distribution. We compute the “Acceptance” ratio :

$$a = \min\left(1, \frac{p(x^*)}{p(x)} \cdot \frac{q(x \rightarrow x^*)}{q(x^* \rightarrow x)}\right)$$

[5]


Gibbs Sampling

like MH but along an axis, useful when we can condition on other variables. Look, we can Gibbs sample in image MRFs with discrete state spaces [2]

Parallel Tempering

(aka “Replica Exchange Monte Carlo” [6], aka “Something to do with your 8 cores”)

- Run N replicas of your chain, each at a different temperature
- Periodically propose MH-style swaps between adjacent chains
- Let's hot chains move around in flatter energy landscape



Swendsen-Wang

Work Through

Data-driven MCMC

- Any “move” is valid as long as it is reversible.
- We can cheat a little bit and construct moves based on the data to help the chains mix, without changing the target distribution

[7] Work Through

Partial Replica Exchange

[4]

MRFs and Parallelism

The conditional independence assumptions allow fine-grained parallelism

Our Implementation

use SW, etc. python, numpy, scipy, c++, boost, etc.
multithreaded C++

How to measure performance? I'm going to go for log-likelihood,

2-D Synthetic Data

3-D Synthetic Data

Div and Audrey

PRELUDE

Why, man, why?

Why go through all this effort for mediocre performance?

- Why be bayesian? - incorporate prior information explicitly - methods compose
- Why use MCMC? - MCMC “black boxes” that compose nicely [?]
- Why use factor graphs? - reasonable prior expressing local dependence - factors allow programmatic evaluation of conditional independencies, thus can easily identify opportunities for parallelism

Where to now?

Exact sampling using Systematic Stochastic Search Better
neighborhood connectivity / likelihood? GPU implementation
Better visualization of posterior?

More information

Try it out

Source and presentation available for download on github

<https://github.com/ericmjonas/mrimrf/>

References I



William Freeman and Egon Pasztor.
Markov networks for low-level vision.
MERL Tech Report, TR99(08), February 1999.



S. Geman and D. Geman.
Stochastic relaxation, Gibbs distributions, and the Bayesian restoration of images, pages 452–472.
Morgan Kaufmann Publishers Inc., 1990.



F.R. Kschischang, B.J. Frey, and H.-A. Loeliger.
Factor graphs and the sum-product algorithm.
Information Theory, IEEE Transactions on, 47(2):498–519,
2001.

References II



Faming Liang and Wing Hung Wong.

Evolutionary monte carlo: Applications to c_p model sampling and change point problem.

STATISTICA SINICA, 10:317—342, 2000.



Nicholas Metropolis, Arianna W. Rosenbluth, Marshall N. Rosenbluth, Augusta H. Teller, and Edward Teller.

Equation of state calculations by fast computing machines.

The Journal of Chemical Physics, 21(6):1087–1092, June 1953.

References III



Robert H. Swendsen and Jian-Sheng Wang.

Replica monte carlo simulation of Spin-Glasses.

Physical Review Letters, 57(21):2607, November 1986.

Copyright (C) 2009 The American Physical Society; Please report any problems to prola@aps.org.



Zhuowen Tu and Song-Chun Zhu.

Image segmentation by data-driven markov chain monte carlo.

Pattern Analysis and Machine Intelligence, IEEE Transactions on, 24(5):657–673, 2002.

References IV



Lei Ying, Zhi-Pei Liang, D.C. Munson, R. Koetter, and B.J. Frey.

Unwrapping of MR phase images using a markov random field model.

Medical Imaging, IEEE Transactions on, 25(1):128–136, 2006.